

Greedy Algorithms




- Solve problems with the simplest possible algorithm
- The hard part: showing that something simple actually works
- Pseudo-definition
 - An algorithm is **Greedy** if it builds its solution by adding elements one at a time using a simple rule

Scheduling Theory

- Tasks
 - Processing requirements, release times, deadlines
- Processors
- Precedence constraints
- Objective function
 - # of jobs scheduled, delay, total execution time

Interval Scheduling

- Given a set of n tasks
- i -th task start and finish time: $s(i)$, $f(i)$
- Two tasks are compatible if they don't overlap
- Goal: find the maximum number of mutually compatible tasks



Greedy template

Let T be the set of tasks, construct a set of independent tasks I

A is the rule determining the greedy algorithm

$I = \{\}$

While (T is not empty)


 Select a task t from T by a rule A

 Add t to I


 Remove t and all tasks incompatible with t from T

Simulate the greedy algorithm for each of these heuristics


A1: Schedule earliest starting task



A2: Schedule shortest available task




A3: Schedule task with fewest conflicting tasks




Greedy solution based on earliest finishing time


Example 1



Example 2



Example 3



EARLIEST-FINISH-TIME-FIRST $(n, s_1, s_2, \dots, s_n, f_1, f_2, \dots, f_n)$

Sort tasks by finish time so that $f_1 \leq f_2 \leq \dots \leq f_n$

$A \leftarrow \Phi$

for $i=1$ to n

 if task i is compatible with A

$A \leftarrow A \cup \{i\}$

Return A

Runtime: $O(n \log n)$ - sorting

To check if task i is compatible with A , just need to keep track of the last added task j^* and compare s_i to f_{j^*}

Analysis of EARLIEST-FINISH-TIME-FIRST

Claim: the EARLIEST-FINISH-TIME-FIRST algorithm is optimal

Key idea: EARLIEST-FINISH-TIME-FIRST stays ahead

- Let $i_1 i_2 \dots i_p$ be the set of tasks selected by greedy ordered by finish time
- Let $j_1 j_2 \dots j_q$ be a set of tasks selected by a different algorithm ordered by finish time

We can show that for $r \leq \min(p, q)$, $f(i_r) \leq f(j_r)$

Stay ahead lemma

Proof (by induction)

Base case:

Greedy choose i_1 with the minimum f , so we must have $f(i_1) \leq f(j_1)$


Induction assumption:

Assume this is true for $f(i_r) \leq f(j_r)$ for $r=1, \dots, k$

Inductive step:


At the end of r -th iteration, task i_{r+1} and j_{r+1} are both compatible with A
Greedy always choose the one with minimum f , thus $f(i_{r+1}) \leq f(j_{r+1})$

QED



Completing the proof

- Let i_1, \dots, i_p be the set of tasks found by EFTF in increasing order of finish times
- Let j_1, \dots, j_q be the set of tasks found by an optimal algorithm in increasing order of finish times
- If $p < q$, then the EFTF stopped before it ran out of tasks - contradiction



- Thus we must have $p=q$