


Minimum Spanning Trees

Adapted from notes by Mahesh Viswanathan (UIUC)

MST: the problem

- Input: Connected graph $G = (V, E)$ with edge costs
- Goal: find $T \subseteq E$ such that (V, T) is connected and the total cost of all edges in T is the smallest
 - T is called the Minimum spanning tree of G



Practical application of MST

- Network design
 - Telephone, computer, road etc.
 - A set of locations and find the cheapest way of connecting them
 - Must be a spanning tree
 - Need to connect all nodes
 - Must be a tree otherwise we can remove some edges and still be connected and cheaper
- Other applications
 - Approximations to NP-hard problems like traveling salesman problem
 - Clustering in machine learning
 - ...

Greedy Template

T is empty (* T will store edges of a MST*)

While T is not spanning tree yet (* $|T| \leq |V| - 1$ *)
 select $e \in E$ to add to T according to a greedy criterion

Return T

How should we choose the edge to add to T ?

We will present two different greedy criteria and then show a unifying proof for the correctness

Kruskal's algorithm

Process edges in the order of their costs (starting from the least) and add edges to T as long as they don't form a cycle.




Figure: Graph G




Figure: MST of G

Kruskal's algorithm

Process edges in the order of their costs (starting from the least) and add edges to T as long as they don't form a cycle.




Figure: Graph G




Figure: MST of G

Kruskal's algorithm

Process edges in the order of their costs (starting from the least) and add edges to T as long as they don't form a cycle.




Figure: Graph G




Figure: MST of G

Kruskal's algorithm

Process edges in the order of their costs (starting from the least) and add edges to T as long as they don't form a cycle.




Figure: Graph G




Figure: MST of G

Kruskal's algorithm

Process edges in the order of their costs (starting from the least) and add edges to T as long as they don't form a cycle.




Figure: Graph G




Figure: MST of G

Kruskal's algorithm

Process edges in the order of their costs (starting from the least) and add edges to T as long as they don't form a cycle.




Figure: Graph G




Figure: MST of G

Prim's algorithm

T maintained by algorithm will be a tree. In each iteration, pick edge with least attachment cost to T .




Figure: Graph G




Figure: MST of G

Prim's algorithm

T maintained by algorithm will be a tree. In each iteration, pick edge with least attachment cost to T .




Figure: Graph G




Figure: MST of G

Prim's algorithm

T maintained by algorithm will be a tree. In each iteration, pick edge with least attachment cost to T .




Figure: Graph G




Figure: MST of G

Prim's algorithm

T maintained by algorithm will be a tree. In each iteration, pick edge with least attachment cost to T .




Figure: Graph G




Figure: MST of G

Prim's algorithm

T maintained by algorithm will be a tree. In each iteration, pick edge with least attachment cost to T .




Figure: Graph G




Figure: MST of G

Prim's algorithm

T maintained by algorithm will be a tree. In each iteration, pick edge with least attachment cost to T .




Figure: Graph G




Figure: MST of G

Prim's algorithm

T maintained by algorithm will be a tree. In each iteration, pick edge with least attachment cost to T .




Figure: Graph G




Figure: MST of G

Correctness

- For now, we will assume that all edge costs are distinct
 - As such there is a unique MST
- This can be relaxed later

Cut property

Cut property in plain English:

If we can partition the graph into two parts, and an edge e is the cheapest edge connecting the two parts, then e must be in the MST.

Now let's talk in Math:

Let $S \subset V$ ($\neq \emptyset$ and $\neq V$). Let $e = (v, w)$ be the minimum cost edge with one end in S and the other end in $V \setminus S$. Then e must be in the MST.

Proof Attempt (by contradiction)

Suppose (for contradiction) that e is not in MST T

- Since T is connected (being a spanning tree), there must be some edge f with one end in S and the other in $V \setminus S$
- Since e is cheaper than f , we can replace f with e to get a cheaper spanning tree
- Cheaper? Yes, but can we be sure it will be a spanning tree?
 - When take an edge of a tree and add another one, it might not be a tree anymore

Error in proof:





Figure: Problematic example. $S = \{1, 2, 7\}$, $e = (7, 3)$, $f = (1, 6)$

Proof of cut property

Proof.



- Suppose $e = (v, w)$ is the cheapest edge connecting two parts S and $V - S$ e.g., $(7, 3)$ for $\{1, 2, 7\}$ vs. $\{3, 4, 5, 6\}$ and it's not in the MST T
- Consider T , there must be a path from v to w
 - because T is connected
- On this path, there must be an edge connecting S and $V - S$
 - In this case it is $(2, 3)$
- If we remove this edge and add e to T , it will be cheaper and will be a tree – so our assumption can't be true, e must be in the MST T

Correctness of Prim's algorithm

- Prim's algorithm: pick the edge with minimum attachment cost to current tree, and add it to current tree.
- Proof of correctness
 - At every iteration, we select the edge e that is the cheapest edge cross the current tree T and the rest of the graph.

Prim's algorithm

T maintained by algorithm will be a tree. In each iteration, pick edge with least attachment cost to T .




Figure: Graph G




Figure: MST of G

Prim's algorithm

T maintained by algorithm will be a tree. In each iteration, pick edge with least attachment cost to T .




Figure: Graph G




Figure: MST of G

Prim's algorithm

T maintained by algorithm will be a tree. In each iteration, pick edge with least attachment cost to T .




Figure: Graph G




Figure: MST of G

Prim's algorithm

T maintained by algorithm will be a tree. In each iteration, pick edge with least attachment cost to T .




Figure: Graph G




Figure: MST of G

Prim's algorithm

T maintained by algorithm will be a tree. In each iteration, pick edge with least attachment cost to T .




Figure: Graph G




Figure: MST of G

Prim's algorithm

T maintained by algorithm will be a tree. In each iteration, pick edge with least attachment cost to T .




Figure: Graph G




Figure: MST of G

Prim's algorithm

T maintained by algorithm will be a tree. In each iteration, pick edge with least attachment cost to T .




Figure: Graph G




Figure: MST of G

Correctness of Kruskal's algorithm

- Kruskal's algorithm starts with every node in its own connected component
- Each iteration it adds the cheapest edge that links two different connected components
- When an edge $e = (v, w)$ is selected, let S be the connected component for v (or w , equivalently)
- e must be the cheapest edge connecting S to the rest of the graph

Kruskal's algorithm

Process edges in the order of their costs (starting from the least) and add edges to T as long as they don't form a cycle.




Figure: Graph G




Figure: MST of G

Kruskal's algorithm

Process edges in the order of their costs (starting from the least) and add edges to T as long as they don't form a cycle.




Figure: Graph G




Figure: MST of G

Kruskal's algorithm

Process edges in the order of their costs (starting from the least) and add edges to T as long as they don't form a cycle.




Figure: Graph G




Figure: MST of G

Kruskal's algorithm

Process edges in the order of their costs (starting from the least) and add edges to T as long as they don't form a cycle.




Figure: Graph G




Figure: MST of G

When edge costs are not distinct

- Order edges lexicographically to break ties
 - Based on this tie breaking, we will have one unique MST
- Both Prim's, kruskal's are optimal with respect to the lexicographical ordering

Implementation & Data structure

Briefly ...

Prim's algorithm

- Idea: Pick edge with minimum attachment cost to current tree, and add to current tree

Implementing Prim's

E is the set of all edges in G
randomly pick a node u_0 and $S = \{u_0\}$
 T is empty (* T stores edge of a MST*)
while $S \neq V$
 pick $e = (v, w) \in E$ such that
 $v \in S, w \notin S$
 e has minimum cost
 $T = T \cup e; S = S \cup w$
return set T

Analysis:

- Number of iterations: $O(|V|)$
- Picking e in each iteration is $O(|E|)$
- Total time $O(|V||E|)$

More efficient implementation

procedure `prim(G, w)`

Input: A connected undirected graph $G = (V, E)$ with edge weights w_e

Output: A minimum spanning tree defined by the array `prev`

for all $u \in V$:

`cost(u) = ∞`

`prev(u) = nil`

Pick any initial node u_0

`cost(u_0) = 0`

$H = \text{makequeue}(V)$ (priority queue, using cost-values as keys)

while H is not empty:

$v = \text{deletemin}(H)$

$S = S \cup \{v\}, T = T \cup (\text{prev}(v), v)$

 for each $\{v, z\} \in E$:

 if `cost(z) > $w(v, z)$:`

`cost(z) = $w(v, z)$`

`prev(z) = v`

`decreasekey(H, z)`

Maintaining the costs for attaching z to S

Analysis: $(O((|V| + |E|)\log |V|))$ - assuming using binary heap for priority queue

- each node is inserted and deleted once from the priority queue ($O(|V| \log |V|)$)
- Each edge is checked once, leading one possible decreasekey ($O(|E| \log |V|)$)

Kruskal's algorithm

- Idea: pick edge of lowest cost and add if it does not form a cycle with existing edges

Implementing Kruskal's

sort the edges in E in increasing order based on cost

T is empty (* T stores edge of a MST*)

for each edge e in sorted order

if $T \cup e$ does not creates cycle

$$T = T \cup e$$

return set T

Analysis:

- Sorting the edges: $O(|E| \log |E|)$
- For loop executes $O(|E|)$ times
- Each time, deciding if adding an edge $e = (u, v)$ leads to cycle:
 - Run DFS or BFS on T to see if u and v are connected $O(|V| + |E|)$
- Total time $O(|E| \log |E| + |E|(|V| + |E|)) = O(|E|^2)$

Efficient Implementation of Kruskal's

- Use **Union-by-rank** data structure to maintain disjoint sets
- Each set contains the nodes of a particular connected component
- Initially each node is in a component by itself

procedure `kruskal` (G, w)

Input: A connected undirected graph $G = (V, E)$ with edge weights w_e

Output: A minimum spanning tree defined by the edges X

for all $u \in V$:

`makeset`(u) ← Place every node in its own connected component. $O(|V|)$

$X = \{\}$

Sort the edges E by weight ← $|E| \log |E|$

for all edges $\{u, v\} \in E$, in increasing order of weight:

 if `find`(u) \neq `find`(v): ← If u and v are not in the same connected component $O(\log |V|)$

 add edge $\{u, v\}$ to X

`union`(u, v) ← Merge the two connected components of u and v . $O(C)$

- Total running time:

$$O(|V|) + O(|E| \log |E|) + O(|E| \log |V|)$$